

# Performance Optimization of the Lattice Boltzmann Method for Computational Fluid Dynamics

Kseniia Ovchinnikova, Anna Grebennikova, Anna Remizova  
UGA

April 25, 2025

- `lbm_comm_init_ex1()` is based on exercise 0, with some changes:
  - Communicator size is  $> 0$ .
  - Total width is divisible by number of processes.
  - X: communicator size; Y remains 1.
- `lbm_comm_ghost_exchange_ex1()`:
  - Computed ranks of left and right neighbors.
  - Neighbors set to `MPI_PROC_NULL` at domain boundaries.
  - Performed column exchange using blocking communication (we first do the yellow passage and wait for it to finish, and then we do the green one).

[illegible]

Figure: Illustration of ghost cell exchange for 2 processes.

# Exercise 2 and 3: Communication Implementations

---

Blocking communication can lead to deadlocks when working with large datasets and many processes. To address this, we implement two alternative strategies:

- **Exercise 2: 1D odd/even communication**

- `lbm_comm_init_ex2()` is based on `ex1`.
- `lbm_comm_ghost_exchange_ex2()`:
  - Even ranks: receive first, then send.
  - Odd ranks: send first, then receive.

- **Exercise 3: 1D non-blocking communication**

- `lbm_comm_init_ex3()` is based on `ex1`.
- `lbm_comm_ghost_exchange_ex3()`:
  - Initialized MPI request array for non-blocking communication.
  - Used `MPI_Isend` and `MPI_Irecv`.
  - Finalized with `MPI_Waitall`.

## Exercise 4: 2D blocking communication, manual copy

---

- `lbm_comm_init_ex4()` is based on `lbm_comm_init_ex1()` with several changes:
  - Initialized 2D grid dimensions with `MPI_Dims_create`.
  - Created a Cartesian communicator with `MPI_Cart_create`.
  - Retrieved process coordinates using `MPI_Cart_coords`.
  - The matrix is stored in column-major order; for up/down exchanges, buffers must be built manually as data is not stored linearly in memory.
  - Allocated `buffer_recv_up`, `buffer_recv_down`, `buffer_send_up`, and `buffer_send_down`.
- `lbm_comm_release_ex4()` was implemented to free the buffers and the communicator.

## Exercise 4: Blocking communication and manual copy

---

- `lbm_comm_ghost_exchange_ex4()`:
  - Defined variables for neighbor ranks.
  - NB: we didn't need to implement a helper function to get the left/right neighbours because we implemented cartesian topology - we simply use `MPI_Cart_shift` for it.
- *Left-right* communication remains the same as before.
- For *up-down* communication we built send buffers manually using `for` loops.
- After communication, we unpacked the received buffers into the mesh.

# Exercise 4: Results

Select exercise 4

```
* Init...
* Splitting: (2 x 2)
* Rank 0: (0, 0)
* Rank 1: (0, 1)
* Rank 2: (1, 0)
* Rank 3: (1, 1)
* Communicate...
* fetch...
* display...
```

000	010	020	030	040	050	060	070	080	090	080	090	100	110	120	130	140	150	160	170
001	011	021	031	041	051	061	071	081	091	081	091	101	111	121	131	141	151	161	171
002	012	022	032	042	052	062	072	082	092	082	092	102	112	122	132	142	152	162	172
003	013	023	033	043	053	063	073	083	093	083	093	103	113	123	133	143	153	163	173
004	014	024	034	044	054	064	074	084	094	084	094	104	114	124	134	144	154	164	174
005	015	025	035	045	055	065	075	085	095	085	095	105	115	125	135	145	155	165	175
006	016	026	036	046	056	066	076	086	096	086	096	106	116	126	136	146	156	166	176
007	017	027	037	047	057	067	077	087	097	087	097	107	117	127	137	147	157	167	177
008	018	028	038	048	058	068	078	088	098	088	098	108	118	128	138	148	158	168	178
009	019	029	039	049	059	069	079	089	099	089	099	109	119	129	139	149	159	169	179
008	018	028	038	048	058	068	078	088	098	088	098	108	118	128	138	148	158	168	178
009	019	029	039	049	059	069	079	089	099	089	099	109	119	129	139	149	159	169	179
010	020	030	040	050	060	070	080	090	100	090	100	110	120	130	140	150	160	170	180
011	021	031	041	051	061	071	081	091	101	091	101	111	121	131	141	151	161	171	181
012	022	032	042	052	062	072	082	092	102	092	102	112	122	132	142	152	162	172	182
013	023	033	043	053	063	073	083	093	103	093	103	113	123	133	143	153	163	173	183
014	024	034	044	054	064	074	084	094	104	094	104	114	124	134	144	154	164	174	184
015	025	035	045	055	065	075	085	095	105	095	105	115	125	135	145	155	165	175	185
016	026	036	046	056	066	076	086	096	106	096	106	116	126	136	146	156	166	176	186
017	027	037	047	057	067	077	087	097	107	097	107	117	127	137	147	157	167	177	187

>>> VALID <<<

## Nota Bene:

Diagonal exchanges are not required. It is sufficient to first exchange *left-right* and then *up-down*. As shown in the figure, corner behavior is handled correctly.

## Exercise 5: 2D Blocking Communication with Derived MPI Datatype

---

- `lbm_comm_init_ex5()`:
  - Reuses `lbm_comm_ghost_exchange_ex4()` and, on top of it, creates an MPI derived datatype (`MPI_Type_vector`) for vertical (non-contiguous) communication.
- `lbm_comm_ghost_exchange_ex5()`:
  - *Left-right* communication remains the same as in Exercise 4.
  - *Up-down* communication simplifies: no need to pack and unpack buffers, we just use the new derived MPI datatype.
- `lbm_comm_release_ex5()`: Reuses `ex4` release logic and on top of it frees the derived MPI datatype.

## Exercise 6: 2D Non-Blocking Communication (also with Derived MPI Datatype)

---

- `lbm_comm_init_ex6()` and `lbm_comm_release_ex6()` reuse `lbm_comm_init_ex5()` and `lbm_comm_release_ex5()` respectively.
- `lbm_comm_ghost_exchange_ex6()` looks very similar to `lbm_comm_ghost_exchange_ex5()` but implements non-blocking communication. We use: `MPI_Isend`, `MPI_Irecv` and `MPI_Waitall`. NB: we use `MPI_Waitall` two times to correctly work with diagonal exchanges. For that we run communication in batches: firstly *left-right* and after it's finished we use *up-down*.



# Results

---

For all the above-described exercises we have done checksums and received desired -:OK in comparison with sequential implementation of Exercise 0. We also admired pretty gifs and debugging results.

---

